

University Course Scheduling Using a Genetic Algorithm

GRADUATE PROJECT TECHNICAL REPORT

Submitted to the Faculty of
the Department of Computing and Mathematical Sciences
Texas A&M University-Corpus Christi
Corpus Christi, Texas

in Partial Fulfillment of the Requirements for the Degree of
Master of Science in Computer Science

by

David D. Barth
Summer 2003

Committee Members

Dr. Michelle Moore
Committee Chairperson

Dr. Dulal Kar
Committee Member

Dr. David Thomas
Committee Member

ABSTRACT

Creating course schedules is a time consuming operation for both humans and computers. When creating a schedule, the scheduler must allocate resources (i.e. rooms, instructors, and meeting times) while ensuring that multiple constraints are satisfied. This paper presents an automated scheduling algorithm that applies the concepts of genetic algorithms to the course scheduling problem. For ease of use, a graphical user interface has been incorporated into the system. The result is an easy to use scheduling algorithm that produces good schedules in a short time frame.

TABLE OF CONTENTS

Abstract	ii
1. Introduction and Background	1
1.1 The Course Scheduling Problem	1
1.2 Genetic Algorithms	2
1.3 Terms of Genetic Algorithms	4
1.4 Using Genetic Algorithms to Solve the Course Scheduling Problem	5
1.5 Background Research	5
2. Course Scheduling System	7
2.1 Data Entry	7
2.2 Schedule Creation	14
2.3 Reports	15
3. System Design	17
3.1 The Database	17
3.2 The Graphical User Interface	19
3.2.1 Database Access	19
3.2.2 Constraint Classes	20
3.2.3 Other Data Entry Classes	21
3.2.4 Integrity	22
3.2.5 Simplicity and Uniformity	25
3.3 The Scheduling Algorithm	25
3.3.1 Representation of Schedules	26
3.3.2 Representation of Constraints	27

3.3.3 Calculating Fitness	28
3.3.4 Selection	29
3.3.5 Recombination	32
3.3.6 Mutation	34
3.3.7 Structure	35
3.3.8 Special Considerations	37
3.3.9 Programming Language	41
4. Evaluation and Results	42
4.1 Testing	42
4.2 Evaluation	43
4.2.1 Running Time	43
4.2.2 Violation Free Schedules	43
4.2.3 Comparison of Simple GA and Steady State GA	44
5. Future Work	48
5.1 Mechanism for Manual Changes	48
5.2 Determination of Data Dependencies	48
5.3 Testing On Larger Schedules	48
5.4 Use of Clusters in Schedule Creation	49
6. Conclusion	50
Bibliography and References	51

LIST OF FIGURES

Figure 2.1 Course Data Entry Form.....	9
Figure 2.2 Courses to Schedule Form.....	10
Figure 2.3 Course Constraints Form Conflicts	11
Figure 2.4 Course Constraints Form (Time).....	12
Figure 2.5 Instructor Constraints Form.....	13
Figure 2.6 Constraint Penalties Form	14
Figure 2.7 Scheduling Algorithm Message	15
Figure 3.1 Entity Relationship Diagram	17
Figure 3.2 Database Class.....	19
Figure 3.3 Base Class.....	20
Figure 3.4 Constraint Classes	21
Figure 3.5 Other Data Entry Classes.....	22
Figure 3.6 Representation of a Schedule	26
Figure 3.7 Representation of a Constraint	28
Figure 3.8 Roulette Wheel Sized According to Fitness.....	30
Figure 3.9 Chromosomes Before Swapping	32
Figure 3.10: Chromosomes After Swapping	33
Figure 3.11 Structure Chart.....	36
Figure 3.12 Bit String Representation of Days.....	38
Figure 4.1 Comparison of sGA and ssGA Over 1000 Generations.....	45
Figure 4.2 Comparison of sGA and ssGA Over 10,000 Generations.....	46
Figure 4.3 Comparison of sGA and ssGA Over 100,000 Generations.....	47

LIST OF TABLES

Table 3.1 Entity Data Dictionary	18
Table 2 Schedule Fitness Values	30

1. INTRODUCTION AND BACKGROUND

1.1 The Course Scheduling Problem

Every semester members of a college's faculty or staff must grapple with the problem of scheduling the courses to be taught in the next semester. Creating these semester schedules is a time consuming and error prone task that causes many frustrations for the person in charge of their creation. When creating a schedule, the scheduler must ensure that every course is assigned a classroom, instructor, and a timeslot.

The difficulty of creating schedules is due mainly to the fact that the scheduler must assign courses to a finite number of resources (e.g. classrooms, instructors, etc). When assigning these resources to a course, the scheduler must be sure that no conflicts are created, such as assigning two courses to the same room at the same time. Scheduling conflicts often go undetected during the planning process and result in incorrect schedules being delivered to students.

Typically, when creating a schedule, the scheduler must contend with a set of constraints. Some colleges deal with the following types of constraints when creating schedules:

1. Some courses may not be taught at the same time (to allow students to take both courses in the same semester). This is the course conflicts constraint.
2. Some courses must be taught at specific times.
3. Some courses must be taught in particular classrooms.
4. A course may be assigned a specific instructor.

5. Instructors may only teach courses for which they are qualified.
6. Instructors have course load requirements (how many hours per week they are willing to spend in the classroom).
7. Instructors have time preferences (instructors may not want to teach at certain times).
8. Instructors cannot teach two courses at the same time.
9. A room cannot have two courses scheduled at the same time.
10. Class size should not exceed the capacity of the scheduled classroom.
11. Some parts of the day may be blacked out, so that courses may not be scheduled during the blackout period.

When creating a schedule, the scheduler must ensure that none of these constraints have been violated, or at least be aware of any violations that were unavoidable. Coping with these constraints is a time consuming process that the scheduler would be more than happy to turn over to a computer.

1.2 Genetic Algorithms

A genetic algorithm (GA) is a problem solving method that is based on the process of natural selection. In nature, a species survives when two individuals of that species reproduce to form an offspring. The process of natural selection chooses the individuals, within a species, that are allowed to reproduce. This selection is made based on the rule of “survival of the fittest”; strong individuals are more likely to reproduce than weaker individuals, which results in stronger offspring.

When two individuals mate, their offspring is a combination of its parents' genetic material. The offspring's chromosomes are created when the chromosomes of its parents interact with one another exchanging genetic information. This interaction is known as recombination.

An individual's genes can also be affected by the process of mutation. A mutation is the permanent alteration of an individual's genes. This change can improve the individual, harm the individual or have no effect on the individual.

A GA operates by simulating the genetic operators of selection, recombination, and mutation. A GA works by creating an initial population of solutions to the problem at hand. Each individual within the population is a complete solution to the problem; therefore a population is made up of multiple solutions. This initial population can be created by some algorithm, or can be created at random.

Once the initial population has been created, those individuals that will be allowed to reproduce are selected based on the individual's fitness value. An individual's fitness is a measure of how well that individual solves the problem. An individual with a higher fitness is more likely to be selected for reproduction than an individual with a lower fitness.

In the recombination phase the individuals selected for mating are paired up and exchange information between one another. The information exchanged between the parents is chosen at random. This information exchange results in the creation of two new individuals that may be a stronger solution than the parents.

The mutation phase involves the changing of information in single individual. The information to be changed is chosen at random, with the hope that making the

change will result in a stronger individual. The individual may also be unaffected or made weaker by performing mutation.

GAs solve problems by performing the following steps:

1. Initialization of the initial population.
2. Evaluation of the fitness of each individual in the population.
3. Selection of the most fit individuals.
4. Changing of the individuals through a process of recombination or mutation.
5. Repeating of steps 2 thru 4 until a desirable solution is found.

These steps are performed regardless of the type of problem being solved, which means that GAs can be applied to a wide range of optimization problems (such as scheduling).

In order to apply a GA to a problem, we must be able to represent a problem solution with some sort of data structure encoding and have a means of determining the fitness of each individual. Section 3 of this report will discuss the encoded representation of individual schedules and evaluation of a schedule's fitness. It will also discuss the manner in which the genetic operators (selection, recombination and mutation) are applied.

1.3 Terms of Genetic Algorithms

The terms of a genetic algorithm are borrowed from the biological study of genetics. An individual is a complete solution to the problem. An individual is made up of chromosomes, where a chromosome is one part of the solution. For example, in a course scheduling system an individual would contain schedules for each of the courses, and a chromosome would be the schedule for a single course. Chromosomes are made up

of genes, which are individual properties of a chromosome such as the time at which a course is scheduled to meet. When a gene takes on a particular value, that value is called an allele.

As the genetic algorithm proceeds, the groups of individuals that make up the search space for the GA change. The individuals that are a member of the search space at any given time are known as a generation.

1.4 Using Genetic Algorithms to Solve the Course Scheduling Problem

The purpose of this project was to develop a system that will create a complete semester schedule without user intervention. In order to create a schedule, the constraints listed in section 1.1 are considered. Algorithms that would yield an optimal schedule are very time consuming and fall into a class of problems known as “NP-complete”. A scheduling algorithm that performs an exhaustive search is NP-complete because as the number of items to schedule increases the execution time increases exponentially. This project uses a genetic algorithm to produce an approximate solution to the course scheduling problem that runs in polynomial time (see section 4.2.1).

1.5 Background Research

Two of the primary concerns of this project were how to represent a schedule of courses and how to perform the genetic operations of crossover and mutation. Dave Corne et al. [Corne 1995] describe a method of representing the individuals and constraints using arrays. This is appealing because it allows for a simple means of storing an individual schedule; it also simplifies the task of constraint checking. This

report will discuss how two-dimensional arrays were used to represent schedules and constraints in a manner similar to the one suggested by Corne. The actual use of two dimensional arrays will be discussed in section 3.

Lars Kragelund [Kragelund 1997] in discussing the solution for a system to schedule doctors to shifts at an emergency room describes different types of GAs. What makes each GA different is the method by which the individuals of a generation are replaced. Kragelund calls the two methods for replacing the individuals of a generation "steady state GA (ssGA)" and "simple GA (sGA)".

With sGA every individual in a population is replaced in the next generation. The replacement occurs through the use of the crossover and mutation operators. ssGA performs a partial replacement of the population between generations. In Kragelund's example individuals are selected for reproduction or replacement through the use of a tournament. A tournament randomly selects a group of individuals. The tournament group is then used to select the two most fit individuals for crossover and the two least fit individuals for replacement. With ssGA, some individuals are able to live for multiple generations, while an individual exist only for one generation with sGA. This project implemented both ssGA and sGA and compared the performance of these methods. The comparison of ssGA and sGA is discussed in section 4.

Kragelund also presents a structure for storing the representations of a schedule. He uses a structure that contains a two dimensional array to contain a schedule and variables to store the fitness value for that schedule.

2. COURSE SCHEDULING SYSTEM

The software created during the course of this project is designed to create complete course schedules without intervention from the user. The scheduling algorithm takes into account the constraints mentioned in section 1.1 and attempts to create a schedule that does not violate any constraints. The scheduling algorithm is designed for use by non-technical users. Because of this, an easy to use GUI is provided and is discussed in this section. For this report, the user will be referred to as the scheduler, the genetic algorithm will be referred to as the scheduling algorithm and the graphical user interface will be referred to as the GUI.

All interaction with the scheduling algorithm occurs through the GUI. The functionality provided by the GUI is divided into three categories:

1. Data Entry
2. Schedule Creation
3. Reports

This section discusses the GUI as it appears to the scheduler. The design and internal operation of the GUI is discussed in section 3.

2.1 Data Entry

The following data is required by the scheduling algorithm in order to create a complete schedule:

1. Course Data.
2. Course Constraint Data
 - a. Conflicts – Some courses may not be taught at the same time.

- b. Instructors – A course must be taught by a specific instructor.
 - c. Times – Some courses must be taught at specific times.
 - d. Classrooms – Some courses must be taught in a particular room.
 - e. Blackout times – Colleges may choose to set aside specific times in which no courses will be taught.
- 3. Instructor Data.
 - 4. Instructor Constraint Data.
 - a. Contact Hour Loads – Instructors can teach a minimum and maximum number of hours per week.
 - b. Courses – Instructors are qualified to teach certain courses.
 - c. Time Preferences – Instructors prefer to teach at certain times.
 - 5. Classroom Data.
 - 6. Building Data – Allows user to specify in which building a classroom is located.

The data entry screens allow the scheduler to enter course information, specify which courses to schedule and apply constraints to those courses. When the program first opens, the user is asked to specify the semester for which input is occurring. All data entry is tied to the selected semester with the exception of:

- 1. Buildings/Rooms
- 2. Courses
- 3. Blackout times
- 4. Instructor qualifications
- 5. Constraint penalty values

Each of the data entry screens allows the scheduler to add, edit and delete data. The methods for performing these actions are uniform across the system. This uniformity reduces the learning time required of the scheduler.

The Course Data Entry form is displayed in Figure 2.1 and allows the scheduler to input information from the university catalog. The course data represents the courses that a university offers; these courses may or may not be taught in a particular semester. Because the course data is not tied to a specific semester, the scheduler must enter the data only once to create schedules for multiple semesters.

The screenshot shows a Windows-style application window titled "Course Data Entry". The window has a blue title bar with standard minimize, maximize, and close buttons. The main area is divided into two sections. On the left, there are four input fields: "Subject" (containing "COSC"), "Course Number" (containing "2420"), "Title" (containing "C Programming II"), and "Contact Hours - number of hours class actually meets per week (not credit hours)" (containing "6"). On the right, there is a list box titled "Courses" containing a scrollable list of course names. The list includes: "BCIS 2390 Systems Analysis & Design", "COSC 1301 Microcomputer Applications", "COSC 1309 Logic Design", "COSC 1420 C Programming I", "COSC 2420 C Programming II" (which is highlighted), "GRPH 1359 Object-Oriented Comp Grap", "IMED 2345 Interactive Multimedia II", "IMED 2415 Web Page Design II", "INEW 2334 Adv Web Page Programmin", "ITCC 1306 Basic Router Configuration", "ITCC 1342 Local Area Mgmt(LAN)", "ITCC 1346 Wide Area Mgmt(WAN)", "ITNW 1325 Fundamentals of Networkin", and "ITNW 2317 Network Security". At the bottom of the window, there are four buttons: "Save", "Delete", "Clear", and "Exit".

Figure 2.1 Course Data Entry Form

This form allows the scheduler to add, modify and delete course data. When the scheduler enters data without selecting from the list the GUI performs an add operation when the **save** button is clicked. When the scheduler selects a course from the list prior to entering data a modify operation (the selected course is modified) is performed when

the **Save** button is clicked. To delete a course the scheduler must first select a course from the list, and then click the **Delete** button, the scheduler is then prompted to verify the delete operation. The **Clear** button, when clicked, empties all of the fields and deselects the selected course. This readies the form for an add operation to be performed. All of the data entry forms in the GUI follow this same method for data entry.

Figure 2.2 displays the Courses to Schedule form which is used to specify which courses are to be considered on the schedule.

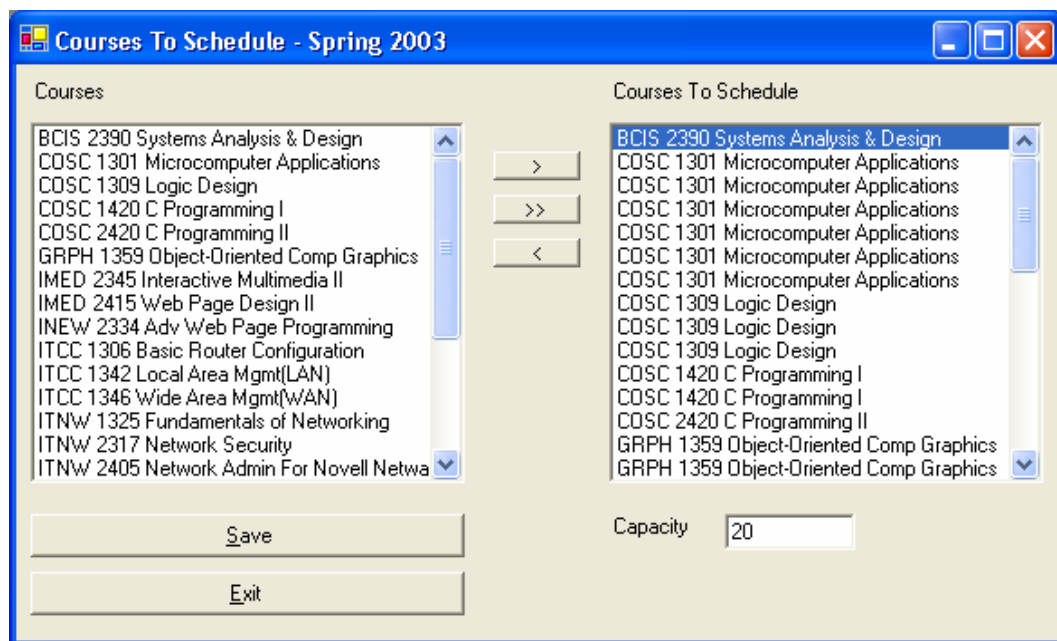


Figure 2.2 Courses to Schedule Form

To add a course to the list of courses to schedule, the scheduler selects a course from the courses list and clicks the **right arrow** button. To remove a course from the schedule the scheduler selects the course from the courses to schedule list and clicks the **left arrow**. Clicking the **double right arrow** moves all courses onto the schedule. The capacity field allows the scheduler to specify the maximum capacity for a course and

defaults to 20 when a value is not specified. The capacity value is associated with an individual course so each course can have a different capacity specified.

The form allows multiple sections of the same course to be specified by simply adding the course the required number of times. When multiple sections are specified, each section is seen as a separate course offering by the GUI and scheduling algorithm. This means that any constraints applied to a particular section of a course are applied to that section and all other sections are unaffected.

Since the courses to schedule are tied to a specific semester, the selected semester is displayed in the title bar (the selected semester is always displayed in the title bar of the main form). If the scheduler selects another semester while this form is open, the form is automatically updated to the new semester. This operation is true for any form that deals with data entry that is tied to a specific semester.

The Course Constraints form, shown in Figure 2.3, allows the user to specify all constraints that are associated with a particular course.

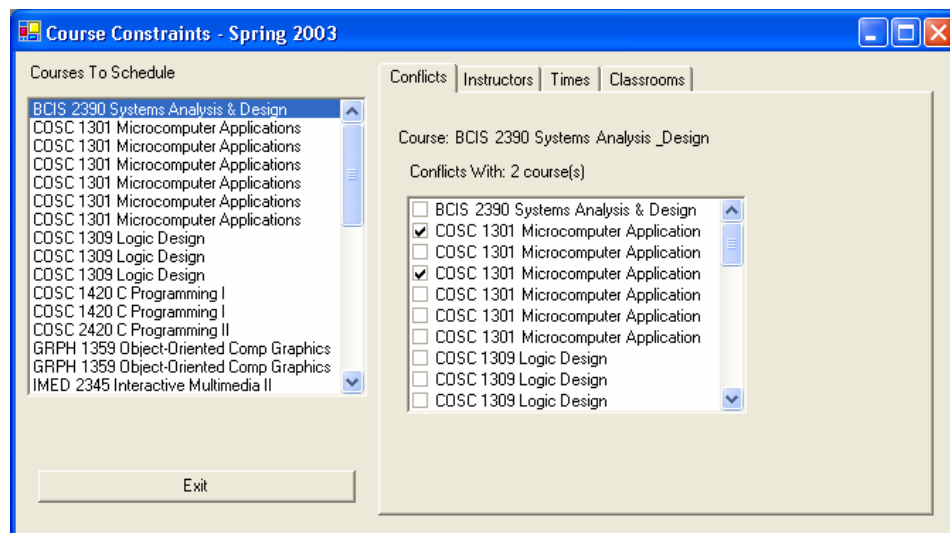


Figure 2.3 Course Constraints Form (Conflicts)

All course constraints are attached to an individual course section. When multiple sections of the same course exist, only the selected section is affected by course constraints.

To operate this form, the scheduler first selects the desired course section from the **Courses to Schedule** list and then applies constraints to the selected section. The four types of course constraints can be accessed by clicking on the appropriate tab. In Figure 2.3 the course constraints form is shown with the **Conflicts** tab selected. To add a new conflict the scheduler simply places a check mark next to the appropriate course. To remove a conflict the scheduler simply removes the check mark. Any constraint data entry that makes use of a checked list follows this method of data entry.

Figure 2.4 displays the Course Constraints form with the Times tab selected. This demonstrates the manner in which days and times are entered by the scheduler. The scheduler simply places a check mark next to each day desired and enters the time in the appropriate box. All times must be entered using a 24-hour format. Any data entry that requires the entry of days and times uses this method.

The screenshot shows a software window titled "Course Constraints - Spring, 2003". It has four tabs: "Conflicts", "Instructors", "Times", and "Classrooms". The "Times" tab is currently selected. On the left, there is a list box titled "Courses To Schedule" containing several course entries, with "BDIS 2390 Systems Analysis & Design" selected. On the right, the "Course: BDIS 2390 Systems Analysis _Design" is displayed. Below this, it says "Is to be taught on: TTH at 930". There are five checkboxes for days of the week: Monday (unchecked), Tuesday (checked), Wednesday (unchecked), Thursday (checked), and Friday (unchecked). To the right of these checkboxes is a text box labeled "Time (Entered in military time, e.g. 0830)" containing the value "930". At the bottom right of the form are "Save" and "Delete" buttons. At the bottom left is an "Exit" button.

Figure 2.4 Course Constraints Form (Time)

A course section can have a single starting time associated with it. After entering the days and times, the scheduler clicks the **Save** button to attribute the days and time to the selected course. An existing day and time can be modified by making the appropriate changes and clicking the **Save** button. The **Delete** button is used to remove the existing days and time from the selected course. Once the days and time have been deleted, the time constraint will not exist for the selected course.

The Instructor Constraints form is shown in Figure 2.5. The basic operations of this form are the same as the Course Constraints form. Care has been taken during the development of the GUI to make the data entry methods as uniform as possible. This uniformity reduces the amount of time required to learn to use the GUI.

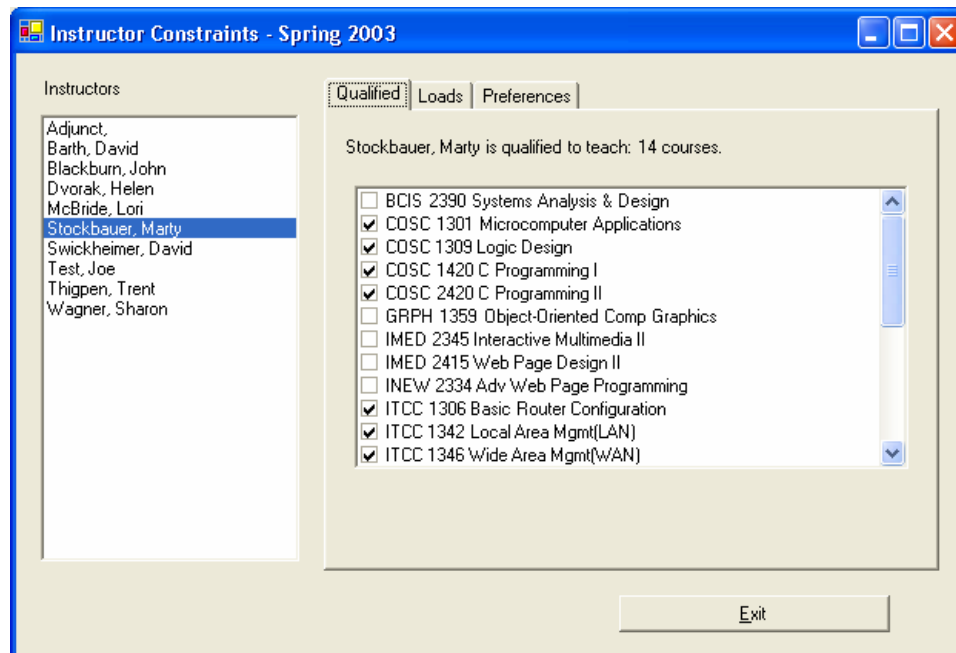
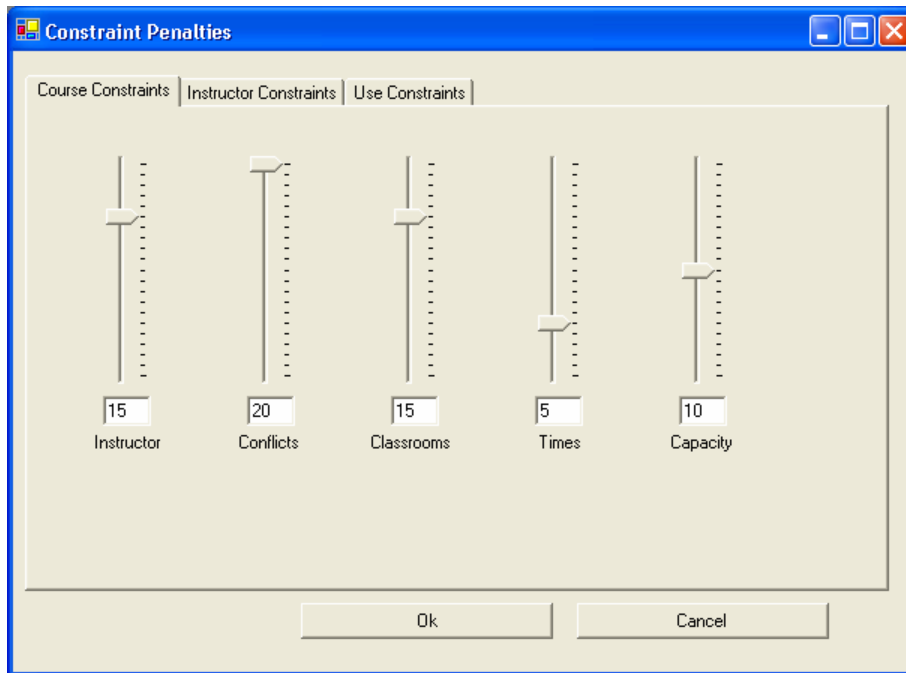


Figure 2.5 Instructor Constraints Form

The scheduling algorithm provides a means of implementing soft and hard constraints. These will be discussed in section 3. The scheduler has the ability to adjust the penalty values of each constraint, thereby determining which constraints are soft and

which constraints are hard. The Constraint Penalties form is displayed in Figure 2.6. It allows the user to adjust the penalty value of each constraint by moving the slider up or down. The user can also change the penalty value by typing the value in the associated box. The constraint penalties can be set to a maximum of 20 and a minimum of 0.



The screenshot shows a window titled "Constraint Penalties" with three tabs: "Course Constraints", "Instructor Constraints", and "Use Constraints". The "Course Constraints" tab is active. It contains five vertical sliders, each with a corresponding input box below it. The sliders are labeled "Instructor", "Conflicts", "Classrooms", "Times", and "Capacity". The input boxes contain the values 15, 20, 15, 5, and 10 respectively. At the bottom of the window are "Ok" and "Cancel" buttons.

Constraint	Penalty Value
Instructor	15
Conflicts	20
Classrooms	15
Times	5
Capacity	10

Figure 2.6 Constraint Penalties Form

2.2 Schedule Creation

Once the scheduler has completed all of the required data entry the scheduling algorithm can be invoked by clicking **Create Schedule** on the schedule menu. The creation of a schedule is completely automated and does not require any user intervention.

The scheduling algorithm runs until a schedule is found that has 5 or fewer constraint violations or until 100,000 generations have elapsed. At that point the user has the choice of accepting the schedule as is or of having the scheduling algorithm continue

running. The scheduler is prompted with the message shown in Figure 2.7, clicking **Yes** causes the scheduling algorithm to continue running while clicking **No** causes the scheduling algorithm to save the schedule and information about the constraint violations to the database.




Figure 2.7 Scheduling Algorithm Message

Once the scheduler has terminated the scheduling algorithm, the completed schedule can be accessed from the reports menu. The scheduler can also view a report listing any constraint violations that exist for the schedule.

2.3 Reports

Reports are included that allow the scheduler to view all data that has been entered into the system and any schedule information created by the scheduling algorithm. The following reports are provided:

1. Data Entry Reports
 - a. Buildings – list all buildings and classrooms in the system.
 - b. Course Constraints
 - i. Blackouts – list all blackout times entered in the system.
 - ii. Classrooms – list all classroom constraints that exist for the given semester.

- iii. Conflicts – list all conflicts that exist for the given semester.
- iv. Instructors – list all courses that must be taught by a specific instructor for the given semester.
- v. Times – list all courses that must be taught at a specific time for the given semester.
- c. Instructor Constraints
 - i. Loads – list each instructor that has a minimum and maximum course load specified.
 - ii. Preferences – list the time preferences entered for each instructor.
 - iii. Qualification – list each instructor along with the courses they are qualified to teach.
- d. Courses – list all courses entered into the system.
- e. Instructors – list all instructors entered into the system.
- 2. Schedule Reports – based on data created by the scheduling algorithm.
 - a. Schedule – list each course along with the instructor, room and time that has been assigned to the course. The report can be sorted by course, instructor, classroom or time.
 - b. Constraint Violations – list all constraint violations that exist on the selected schedule.

A sample of each of the reports listed is contained in Appendix A.

3. SYSTEM DESIGN

The Course Scheduling System consists of three parts: a GUI to allow user interaction, the scheduling algorithm that creates the actual schedules, and a database. Each of these components will be discussed separately in this section.

3.1 The Database

A Microsoft Access database was used to handle the data storage requirements of course scheduling. These requirements include the storage of user input, such as course and constraint information, as well as the storage of completed schedules.

Figure 3.1 is an ER diagram representing the database schema.

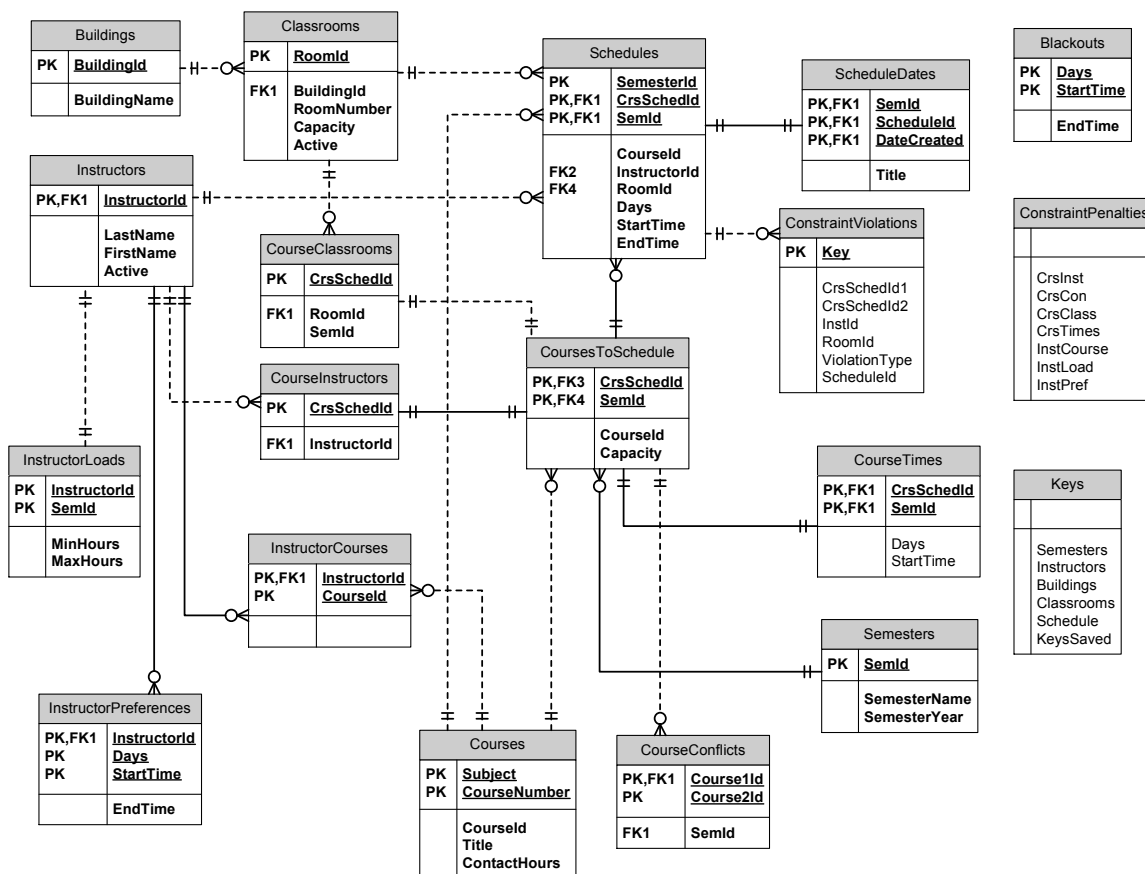


Figure 3.1 Entity Relationship Diagram

The database schema has been designed to keep all of the data needed to create and store a course schedule. The *Blackouts*, *ConstraintPenalties* and *Keys* entities contain setup information used by the scheduling algorithm, and as such are not related to any other entity. The data dictionary shown in Table 3.1 defines each of the entities. The data dictionary in Appendix B defines the attributes of each entity.

Table 3.1 Entity Data Dictionary

Entity	Description
Buildings	Contains a list of buildings used by the department. This data is used for display purposes in the GUI, but is not used by the scheduler.
Blackouts	Blackouts allow the user to specify times of day that the university does not offer any courses.
Classrooms	List the classrooms that a building contains.
ConstraintPenalties	Contains the numeric penalty for each constraint violation.
ConstraintViolations	Contains information about all of the constraint violations contained on a schedule.
Courses	Contains a listing of the courses that may be scheduled.
CourseClassrooms	Courses that must be taught in specific classrooms are listed in the CourseClassrooms entity.
CourseConflicts	Courses that can not be taught at the same time are entered into the CourseConflicts entity.
CourseInstructors	Some courses must be taught by a specific instructor.
CourseTimes	Stores information for courses that must be taught at a specific time.
CoursesToSchedule	Stores the courses that are to be scheduled for the specified semester.
Instructors	A listing of all instructors that should be considered by the scheduler.
InstructorCourses	Records the courses that an instructor is qualified to teach.
InstructorLoads	Specifies the maximum number of hours that an instructor should spend in the classroom each week.
InstructorPreferences	List the times of day that an instructor prefers to teach.
Keys	Stores the primary key values for database tables. Allows the GUI to assign the primary key value.
Schedule	Used to store schedule information for each course. This information is generated by the scheduler, and is not from user input.
ScheduleDates	Stores the date and user given name for each schedule. Allows the user a means of distinguishing between multiple schedules.
Semesters	List semester information that is stored so that the system can store schedules for more than one semester.

The *Schedules* and *ConstraintViolations* entities receive their information from the scheduling algorithm and are read only for the scheduler.

3.2 The Graphical User Interface

The GUI is designed to give the scheduler an easy to use, uniform means of interacting with the database and the scheduling algorithm. The GUI was implemented using Microsoft Visual Basic .Net (VB .Net).

3.2.1 Database Access

GUI access to the database was provided through the use of the Microsoft Active Data Object .Net (ADO .Net). The methods needed for database interaction is encapsulated by the database class (see figure 3.2). When the GUI is first started, a global database object is created. This object is then used by every portion of the GUI that needs database access. The advantage here is that only a single database object is created for the entire operation of the GUI.

Database
-dbConnection : Object
+open() +close() +getResultset() : Object +executeQuery() +getDataStream() : Object

Figure 3.2 Database Class

Once the database connection has been established, it can be controlled through the *open* and *close* methods. The *getResultset*, *executeQuery* and *getDataStream* methods receive structured query language statements and use the database connection to execute the queries. The *executeQuery* method is used to execute action queries and as such does not return a value. The *getResultset* and *getDataStream* methods are used to

execute select queries, *getResultset* returns a data adapter object and *getDataStream* returns a data reader object.

3.2.2 Constraint Classes

Each of the constraints, with the exception of the instructor and room use constraints, is represented in the GUI with a class. These classes provide the ability to retrieve constraint information from the database and the ability to add, modify and delete constraint records.

Each constraint class inherits the same base class that provides basic functionality. The base class includes the objects needed to interact with the database (via the database class), and provides the ability to delete, modify, find and save constraint records. Because adding records varies for each constraint type, the base class does not include an add method. Figure 3.3 displays the attributes and methods of the base class.

clsBase
#dbAdapter : Object
#dbDataSet : Object
#dbCurrRow : Object
#popAdapter()
+find() : Boolean
+delete() : boolean(idl)
+save() : Boolean

Figure 3.3 Base Class

The base class attributes *dbAdapter*, *dbDataSet* and *dbCurrRow* are declared with protected access to allow the derived classes to use them directly. For this same reason, the *popAdapter* method is also declared with protected access. The inheritance relationship between the constraint violations classes and the base class is shown in figure 3.4.

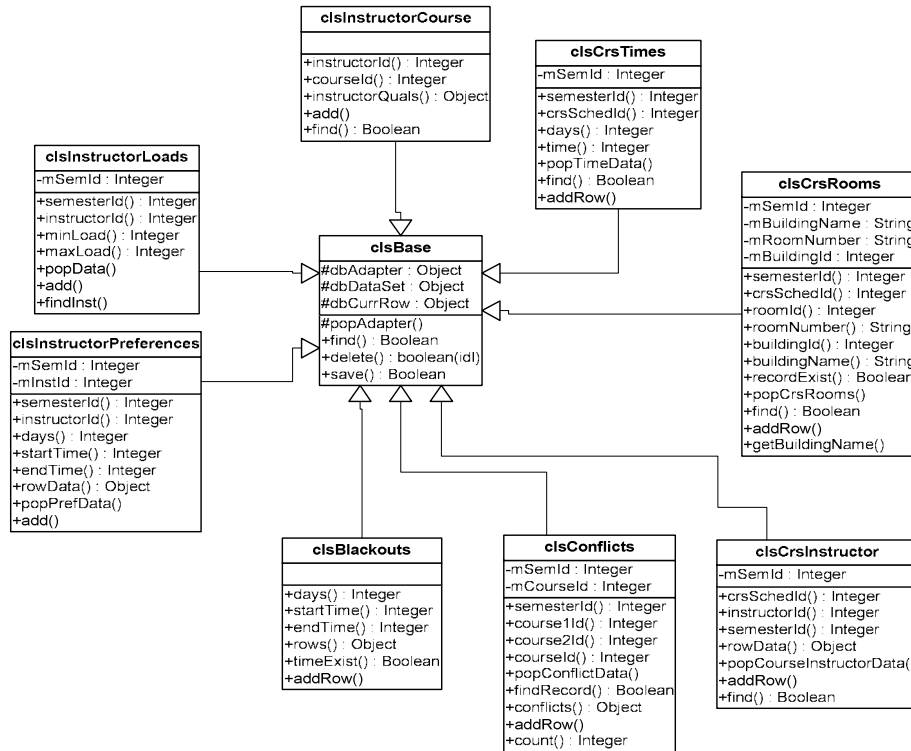


Figure 3.4 Constraint Classes

The room use (a room cannot have two courses scheduled at the same time) and instructor use (an instructor cannot teach two courses at the same time) are implicit constraints that are not specified by the scheduler. Because of this, there is not a class that represents these constraints.

3.2.3 Other Data Entry Classes

Classes are provided to allow for the data entry of buildings, instructors, rooms, courses, semesters and constraint penalty values. Like the constraint classes, these classes provide the ability to retrieve information from the database, and the ability to add, modify and delete information. These classes, like the constraint classes, inherit *clsBase* for their basic database functionality. Figure 3.5 shows the inheritance relationship of these classes.

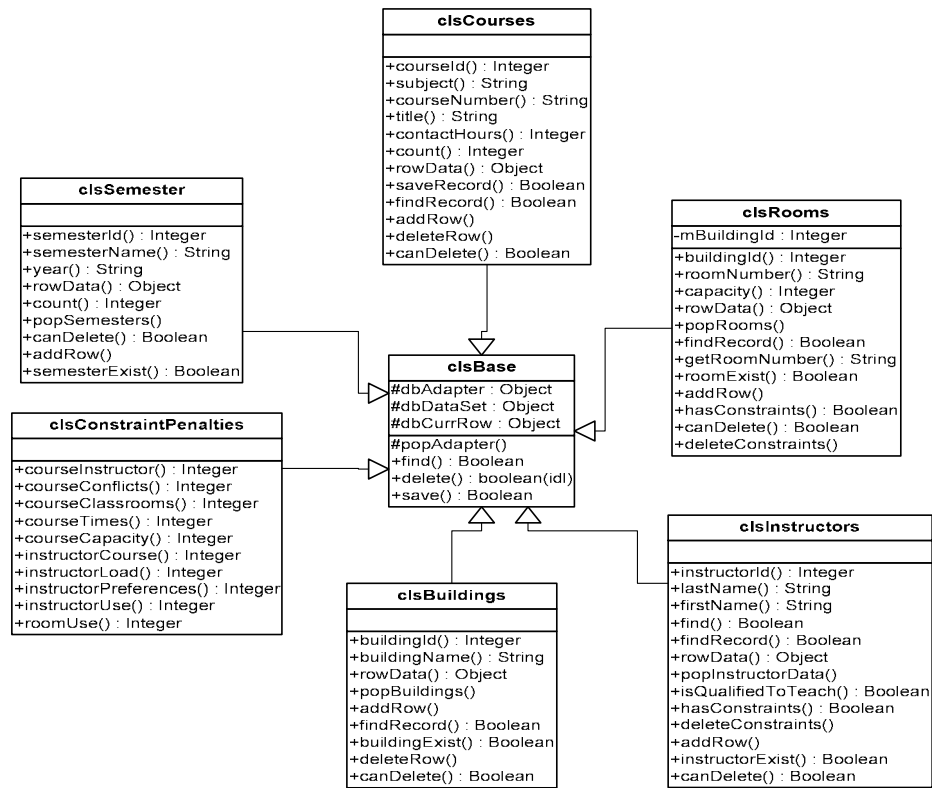


Figure 3.5 Other Data Entry Classes

3.2.4 Integrity

The scheduling algorithm is dependent on the correctness of its input data. If the data provided to the scheduling algorithm is incorrect then it will not be possible to arrive at a correct schedule. To provide the scheduling algorithm with the correct data the GUI must enforce referential integrity and data integrity within the database.

Referential Integrity

Referential integrity is the process of ensuring that a child record in a table always has a corresponding parent record in the parent table. If for example, a course is deleted but constraints associated with that course are left in the database, those constraints would be considered by the scheduling algorithm. These orphaned constraints may

prevent a good schedule from being reached. Because of this, referential integrity is very important to the scheduling algorithm.

Referential integrity can be enforced by the database engine, however in this project it was decided to enforce referential integrity within the GUI. This decision was made because of the increased flexibility in providing error messages to the scheduler and in program flow (error handling does not have to be implemented to handle referential integrity).

When the scheduler attempts to delete a record, the GUI queries any related tables looking for child records. If a child record is found in any related table, then the scheduler is given the option of deleting all child records. If the scheduler declines to delete the child records the parent record deletion is denied.

Data Integrity

The data integrity enforced by the GUI is based on the data requirements that are defined by the scheduling algorithm. The GUI must ensure that the scheduler, when performing data entry, does not input constraints that cannot be fulfilled. Data integrity is crucial because the errors that can be made are subtle and may not be readily apparent to the scheduler. For example, if the scheduler were to apply a constraint specifying the instructor for a course and the instructor was not qualified to teach the course a data error would exist. The scheduling algorithm would not be able to make an instructor assignment that would not violate one of the two constraints.

Whenever a constraint is input, the GUI checks to see if it conflicts with any other constraints already entered into the system. If a conflicting constraint exists then the

constraint being input is rejected. The GUI checks for the following conflicting constraints:

1. A course cannot be set to conflict with itself.
2. Instructors cannot be specified for a course they are not qualified to teach.
3. A course time cannot be set during the school's blackout times.
4. An instructor cannot be given a time preference that occurs during the school's blackout times.
5. A room cannot be specified for a course if the room's capacity is less than the course's capacity.

By enforcing the above rules, the GUI can ensure that conflicting constraints cannot be entered.

Class Days

When assigning meeting times to a course, the scheduling algorithm will only assign days that are combinations of Monday, Wednesday and Friday or combinations of Tuesday and Thursday. For example, a course could be assigned meeting days of Monday, Monday, Wednesday or Monday, Wednesday and Friday. This is done to prevent the assignment of meeting days that would not be desirable at many institutions.

When the scheduler specifies a meeting time for a course, the GUI must ensure that the days entered for the course meet the requirement described above. Any meeting time that does not meet the requirement is rejected by the GUI.

3.2.5 Simplicity and Uniformity

The course scheduling system was designed for non-technical users who may have a very limited knowledge of computer operation. Because of this, every attempt has been made to produce a GUI that is easy to use and that is uniform in operation. The uniformity is extremely important, because as stated earlier uniformity shortens the amount of time required to learn how to use the system.

The operation of the scheduling algorithm is completely hidden from the scheduler's view. This is an important feature that significantly contributes to the simplicity of the system. The scheduler does not have to have any knowledge of how a schedule is created to use this system, but must only understand which constraints can be applied and how to perform the data entry.

3.3 The Scheduling Algorithm

In order to apply a genetic algorithm to the scheduling problem several factors must be considered. These factors are:

1. How will a schedule be represented?
2. How will a constraint be represented?
3. How will a schedule's fitness be determined?
4. How will the selection process be implemented?
5. How will recombination be implemented?
6. How will mutation be implemented?

This section will discuss each of these questions and will also discuss special considerations that arose during the implementation of the scheduling algorithm.

3.3.1 Representation of Schedules

Schedules are represented with a two dimensional integer array. Each row of the array represents a single course and each element of the row represents an assignment that has been made to that course.

	0	1	2	3	4	5	6	7	8
	Course Schedule Id	Course Id	Room Id	Day Assignment	Begin Time	End Time	Instructor Id	Contact Hours	Capacity
0	1	23	2	5	800	930	3	3	25
1	2	34	4	21	1030	1230	5	6	20

Figure 3.6 Representation of a Schedule

Figure 3.6 shows the manner in which a two dimensional array of integers is used to represent a schedule. The heading in each column indicates the values stored in that column, while the numbers on the far left represent the row index and the numbers across the top represent the column index. The *CourseScheduleId* and *CourseId* are the primary keys of the *CoursesToSchedule* and *Courses* tables respectively. They are included in the schedule array for the purpose of comparing constraints to the schedule and saving the completed schedule to the database. The contact hours and capacity are included for the purpose of constraint checking. The *room id*, *day assignment*, *beginning time*, *ending time*, and *instructor id* are the only values assigned by the scheduling algorithm.

The following structure was created to represent a schedule.

```
struct schedule
{
    int (*sch)[9];
    float fitnessValue;
    int totalViolations;
};
```

The member *sch* is a pointer to the two dimensional array which must be dynamically allocated since the number of courses to schedule is not known at design time. The members *fitnessValue* and *totalViolations* allows each schedule to contain its own fitness and number of constraint violations.

The scheduling algorithm creates 40 schedules assigning values to the *room id*, *instructor id*, *days*, *start time*, and *end time* fields at random. These schedules are contained in an array of schedule structures.

3.3.2 Representation of Constraints

Constraints, like schedules, are represented using a two dimensional array of integers. The constraint arrays are also dynamically allocated since the number of constraints is not known at design time. The number of columns in the array is dependent upon the type of constraint being represented. Since the syntax of C++ requires the number of columns to be specified at design time, the columns are set to the maximum number of columns needed (not all constraints require 4 columns). Constraint arrays are created using the structure shown below:

```
struct constraint
{
    int (*cons)[4];
    int penalty;
};
```

The *penalty* field allows the scheduler to set the penalty value for each constraint based on specific schedule preferences. When the scheduling algorithm creates the

constraint arrays, it queries the database for the penalty value for each constraint. Figure 3.7 illustrates the use of a two dimensional array to represent the course conflicts constraint.

Course 1 Id	Course 2 Id
1	2
1	3
3	4

Figure 3.7 Representation of a Constraint

It can be seen in this array that course 1 cannot be taught at the same time as course 2 and 3, and that course 3 cannot be taught at the same time as course 4.

3.3.3 Calculating Fitness

A schedule's fitness is determined by comparing each course in the schedule to each of the constraints. Whenever a constraint violation is found, the constraints penalty value is added to the schedule's total fitness. Once all constraints have been compared, the schedule's fitness value is set equal to $1 / \text{total fitness}$. The division is performed so that the best schedule will have the highest fitness value.

```
int courseInstructorFitness(struct schedule &s)
{
    int instructorFitness = 0;
    int index;
    int crsIndex = 0;

    for(index = 0; index < totalCrsInstCons; index++)
    {
        while(s.sch[crsIndex][0] != crsInst.cons[index][0] &&
            crsIndex < totalCourses)
            crsIndex++;

        if(crsIndex < totalCourses)
        {
            //if a course is not taught by the correct
            //instructor
            if(s.sch[crsIndex][6] != crsInst.cons[index][1])
            {
                //then a constraint violaion exist
                instructorFitness += crsInst.penalty;
                s.totalViolations++;
            }
        }
    }
    return instructorFitness;}

```

The function above is used to compare a schedule against the course instructor constraint. The data type of s is the structure presented in subsection 3.3.1 and the data type of $crsInst$ is the structure presented in subsection 3.3.2. The outer for loop is used to process each of the course instructor constraints. The inner while loop searches through the schedule until the course specified by the current constraint is found. Once the course has been found the assigned instructor is compared to the instructor specified by the constraint. If these values do not match, then a constraint violation exists. When a constraint violation has been determined the constraint's penalty value is added to the total and the schedule's count of violations is incremented. Each of the constraint checking functions is implemented in a manner similar to that shown above.

3.3.4 Selection

Selection is the process of determining which individuals in a population will be allowed to reproduce and create offspring in the next generation. When two individuals reproduce, they are replaced in the next generation by their offspring. This project implemented two versions of schedule replacement: simple GA (sGA) and steady state GA (ssGA). In section 4 the performance of sGA and ssGA will be compared.

The type of replacement scheme used affects the type of selection scheme that can be used. With sGA all of the members of the current generation are replaced by new individuals in the succeeding generation, so the number of schedules selected must be equal to the number of schedules in the generation.

For sGA, selection was implemented by using a method described by Goldberg [Goldberg 1989]. This method uses a roulette wheel to determine which individuals will be selected for reproduction. The wheel is biased in that each individual in the

population “has a roulette wheel slot sized in proportion to its fitness” [Goldberg 1989].

Consider this example, which is based on an example found in [Goldberg 1989].

Suppose that we have 4 complete schedules with fitness evaluations shown in Table 2.

Table 2 Schedule Fitness Values

Schedule	Penalty	Fitness	% of Total
1	300	0.0033	25
2	170	0.0058	44
3	550	0.0018	13
4	400	0.0025	18
Total		0.0131	100

In this example, the penalty value corresponds to the total number of constraint violations found in the schedule. This means that the best schedule has the lowest penalty value. In order to assign the largest fitness value to the best schedule, the fitness value is defined as $1 / \text{penalty}$.

Once the fitness values have been assigned, the roulette wheel is created by giving each schedule a percentage of the wheel equal to its percentage of the total fitness value (see Figure 2.13). Schedules are chosen for reproduction by spinning the wheel (4 times in this case). Schedule 2 was given the largest percentage of the wheel since it has the best fitness. This means that schedule 2 is more likely to be chosen for reproduction than are the other schedules.




Figure 3.8 Roulette Wheel Sized According to Fitness

In this example we would choose 4 schedules for reproduction with the possibility of choosing the same schedule more than once. Since the roulette wheel is based on the fitness values of the schedule, the fitter schedules are favored.

The roulette wheel is represented as an array of integers with 100 elements. Each element in the array stores the location (index) of the schedule (in the array of schedules) assigned the slot. The number of array elements in a slot is determined by the schedule's percentage of the total fitness. For example, if a schedule's fitness was 15% of the total fitness (the sum of the fitness of all schedules) then the schedule would have 15 array elements in its slot.

With ssGA there is only a partial replacement of individuals from one generation to the next. In this project ssGA was implemented using a method described by Kragelund [Kragelund 1997]. The selection process for ssGA is performed by selecting two parent schedules that will be allowed to reproduce and selecting two schedules that will be replaced.

The parent schedules are chosen through the use of a tournament. A tournament selects the two most fit individuals from a randomly selected group of individuals. The system then selects the two least fit individuals from the generation. In the next generation, the two least fit individuals are replaced by the offspring of the two most fit individuals. The most important difference between sGA and ssGA is that with ssGA, some individuals are able to live for multiple generations, while an individual exist only for one generation with sGA.

3.3.5 Recombination

The recombination operator, also known as crossover, creates the new individuals in the next generation. The individuals that were selected by the reproduction operator are used in crossover; these individuals make up the mating pool. Individuals in the mating pool are paired with a mate at random.

Recall from section 1.2.1 that an individual is a complete solution to the problem, in this case a complete schedule, and that an individual is made up of chromosomes. In this case the chromosome represents the assignments (days, time, room, and instructor) that have been made to a particular course. In a course schedule, the chromosomes are represented by the rows of the two dimensional array (see section 3.3.1). Each individual element of a row represents a gene, where a gene describes how the schedule looks. For example, in a schedule a gene represents an instructor assignment . The value of this gene, known as an allele, describes which instructor is teaching the course.

The crossover operator exchanges genes between the chromosomes of the mating pair of individuals. For example, consider the chromosomes in Figure 3.9.




Figure 9 Chromosomes Before Swapping

Each pair of genes has a 50% probability of being swapped. If a pair of genes is chosen, the crossover is achieved by exchanging values between the genes. If the first and last

chromosomes are chosen, then after the swap the resultant chromosomes are seen in Figure 3.10.




Figure 3.10: Chromosomes After Swapping.

The sGA and ssGA both use this method of crossover, commonly called "uniform crossover". The only difference between the two is in the manner in which individuals are chosen for replacement. Every member of a generation is replaced in the succeeding generation with sGA, while only a subset of the individuals are replaced in the succeeding generation with ssGA.

Pairing Chromosomes

Crossover is performed by swapping gene values between the chromosomes of two individuals. The scheduling algorithm performs crossover by swapping gene values between like chromosomes. In other words, the chromosome representing *course1* in a schedule always exchanges genes with the chromosome representing *course1* in the other schedule.

Swapping Class Meeting Times

As seen in figure 3.6 the days a class meets, the starting time, and ending time are each separate genes in a chromosome. When the scheduling algorithm assigns a meeting time to a course, it assigns the meeting days and start time at random. The ending time is assigned by adding the number of hours the course must meet per day (contact hours divided by the number of days the class meets) to the starting time. When swapping gene

values, these three genes must be swapped as a unit, otherwise the meeting times created will not be valid.

3.3.6 Mutation

Based on the course and constraint data entered by the user, there are a finite number of valid schedules that can be generated using that data. This set of possible solutions is the problem's search space.

Since the initial population of solutions is generated at random, it is possible that the gene values generated will not allow the algorithm to move through the entire search space. It is also possible that in the process of swapping gene values, that the reproduction and crossover operators may lose the combination of gene values that are needed to move throughout the search space.

The mutation operator attempts to insure that the individuals have gene values that allow the entire search space to be explored. Mutation works by changing, at random, the value of an individual gene within a chromosome. Mutation is used sparingly, because too much mutation will reduce the effects of reproduction and crossover and cause the individuals to become less fit. In this project mutation has a 5% chance of occurring.

Mutation is implemented by giving each chromosome a 5% chance of being mutated. When a chromosome is mutated, a single gene value (chosen at random) is changed. If any of the genes that make up a course's meeting time is chosen for mutation, the entire meeting time (days, start time, and end time) are mutated. This is done for the same reasons discussed in the section on crossover.

3.3.7 Structure

The scheduling algorithm has been decomposed into multiple modules. The structure chart in figure 3.11 shows this decomposition.

The *Init Algorithm* module retrieves initial data from the database and invokes the creation of the constraint and schedule arrays. Each module that creates a constraint array queries the database for the constraints and transfers them into a two dimensional array. When the schedule arrays are created, course information is retrieved from the database and stored into the two dimensional arrays (multiple schedule arrays are created). Assignments are made to each course in the schedule by the *Make Room Assignment*, *Make Instructor Assignment*, and *Assign Days/Time* modules. These assignments are made at random, but the system ensures that invalid assignments are not made.

The *Calculate Fitness* module calculates the fitness of each schedule in the population, by invoking the module for each type of constraint. Once the fitness of each schedule has been calculated, the roulette wheel is allocated by the *Create Wheel* module. The *Create Wheel* module is not used for the ssGA because ssGA uses a tournament to select schedules for reproduction.

Mutation is carried out by the *Mutate* module and has a 5% chance of occurring. When mutation does occur a gene is chosen at random and given a new value by the *Make Room Assignment*, *Make Instructor Assignment*, or *Assign Days/Times* module. These are the same modules used in the initial schedule creation.

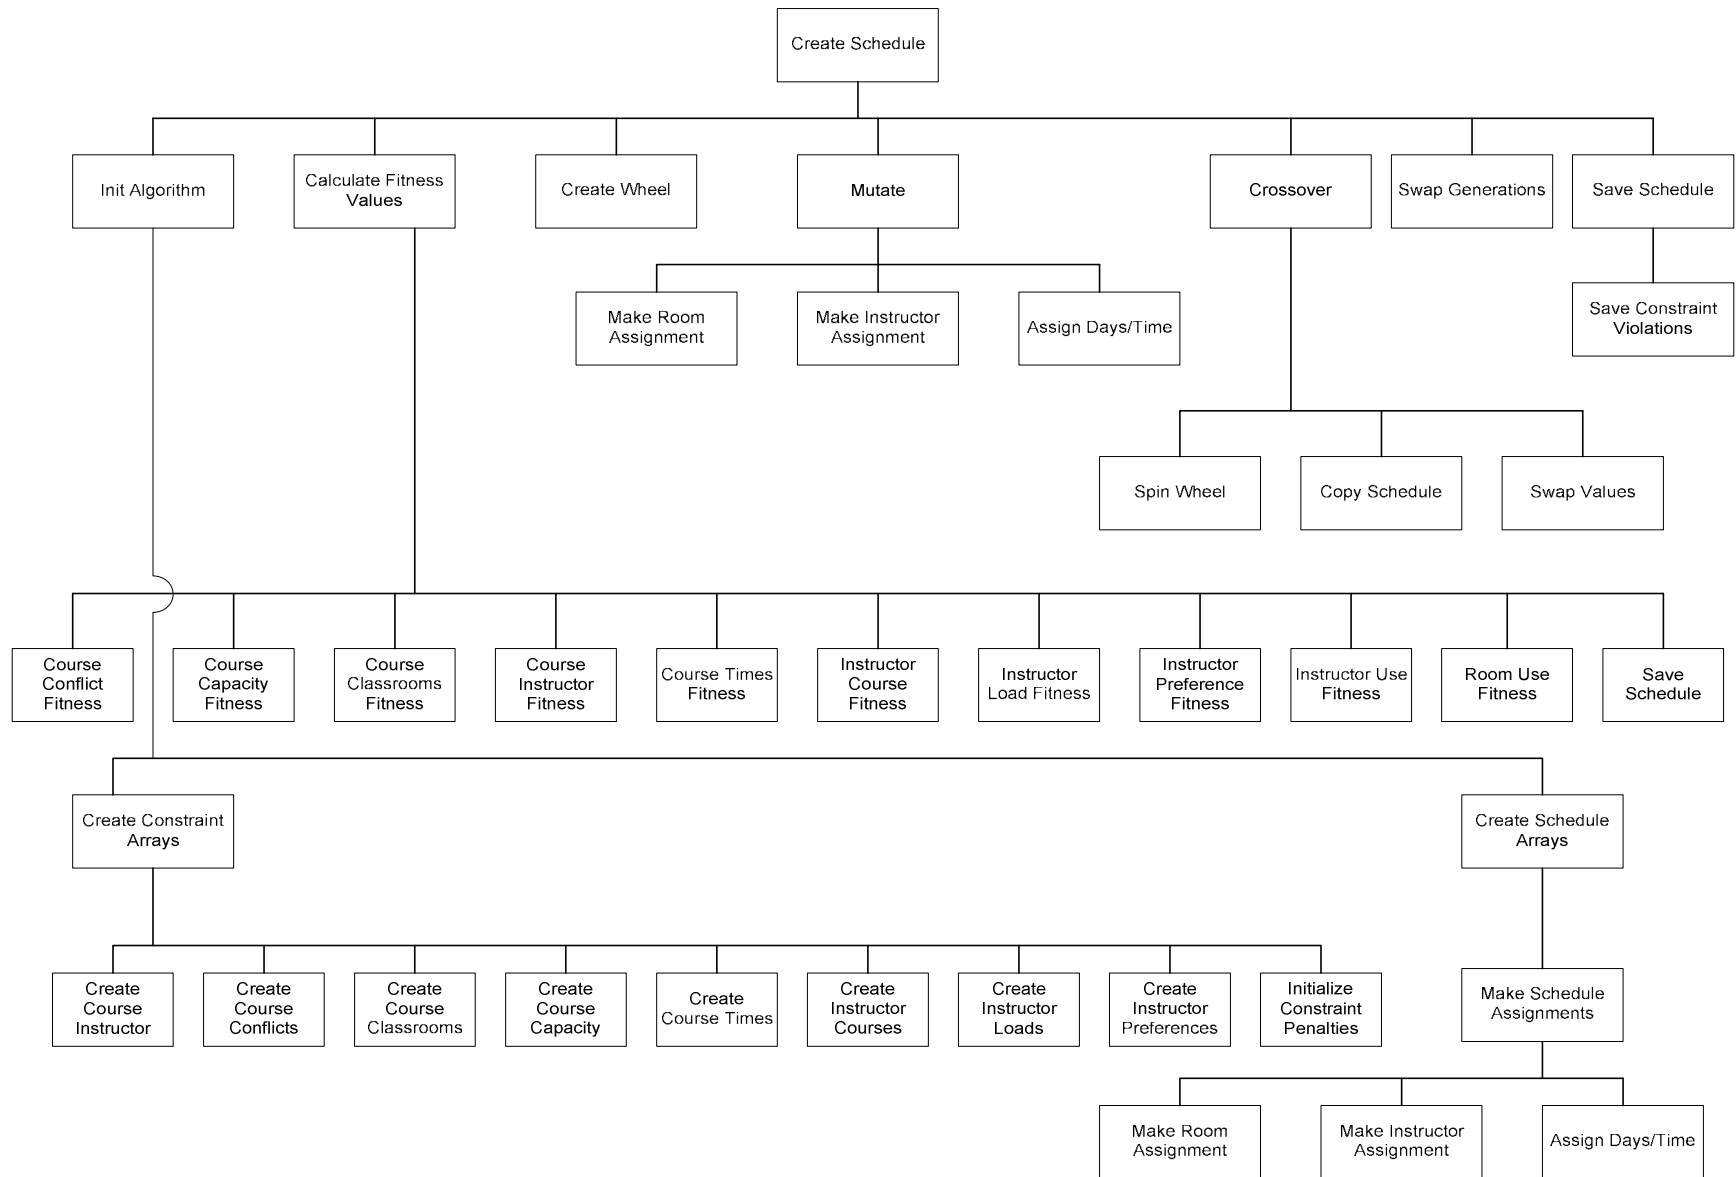


Figure 3.11 Structure Chart

Crossover occurs every generation and operates by first invoking the *Spin Wheel* module (a tournament is used in the case of ssGA) to select schedules for reproduction. Once the schedules have been selected, they are copied into the next generation by the *Copy Schedule* module. The *Swap Values* module exchanges gene values between the schedules in the manner described in section 3.3.5.

When the genetic algorithm has completed running, the *Save Schedule* module is invoked to save the created schedule. The operations of the scheduling algorithm produce a single schedule that is saved to the database. This schedule is the best schedule found across all generations of schedules.

The structure chart shown above does not include all modules that exist in the scheduling algorithm. For clarity some auxiliary modules have not been included, however the structure chart captures the main structure of the scheduling algorithm and provides an accurate depiction of its inner workings.

3.3.8 Special Considerations

During the implementation of the scheduling algorithm several concerns became evident and had to be addressed. These concerns are discussed in this section.

Population Size

The scheduling algorithm uses a small population of schedules in each generation. A small population is used since the scheduling algorithm is designed to run on a desktop computer which may have limited resources. The population size is currently set at 40 schedules. This number was determined through trial and error and seems to yield good results.

Representation of Days and Times

In order for the scheduler to identify a specific time for a course, the days of the week on which the course meets must be represented in some form. This is a problem because a class can meet on any combination of days during a week. To address this problem, courses are represented using a 5 digit binary string where the most significant bit represents Friday and the least significant bit represents Monday.

Representing the days a course meets in this manner simplifies the comparison between two courses when checking constraints such as “Some courses may not be taught at the same time.” For example consider two courses that meet on MWF, the bit string representation of these days would be 10101. These can easily be compared to another course by performing a bit wise AND operation. Suppose that one course meets on MWF and another course meets on MW, to determine that they meet on common days we would perform the AND operation:

$$\begin{array}{r} \phantom{\text{AND}} 1 \ 0 \ 1 \ 0 \ 1 \\ \text{AND} \ 0 \ 0 \ 1 \ 0 \ 1 \\ \hline 0 \ 0 \ 1 \ 0 \ 1 \end{array}$$

Figure 3.12 Bit String Representation of Days

If the AND operation yields a binary string not equal to zero then the courses meet on a common day.

Times are represented using a 24 hour clock. This is also done to simplify processing. By using a 24 hour clock, the system does not have to be concerned with whether a time specifies AM or PM when comparing two times. A course that is

scheduled to meet in room 100 on MWF from 09:00 to 10:00 would have a conflict if another course had been scheduled to meet in room 100 on MWF from 08:00 to 10:00. The 24 hour clock representation of times allows this conflict to be easily detected by the test:

```
if (Course2.beginTime <= Course1.endTime &&  
    Course2.beginTime >= Course1.beginTime)  
    //A time conflict exist
```

Assigning Meeting Times

As discussed in section 2, the days a class meets is restricted to combinations of Monday – Wednesday – Friday and Tuesday – Thursday. This was done to prevent the assigning of non-traditional meeting days and to limit the possibility of courses conflicting with one another.

The blackout constraint is enforced when times are assigned to the courses. The system will not assign a time that violates a blackout constraint. This prevents invalid data from entering the schedules and will allow the process of crossover to be more effective.

Assigning of Instructors and Rooms

When assigning instructors to a schedule, the scheduling algorithm simply generates at random an integer value that corresponds to the instructor id (primary key) value in the database. Since these instructor id values are not sequential (since the user can delete instructors all primary key values will not be present) the system must ensure that all instructor ids assigned are valid. This same problem holds true for the assignment of classrooms.

Soft and Hard Constraints

Since the scheduling algorithm is usually not able to arrive at schedules that are completely free of constraint violations (discussed in section 4) soft and hard constraints have been introduced. A soft constraint is a constraint whose violation is acceptable. Said another way a soft constraint is a preference, we prefer for the constraint to be met, but if it is violated the schedule will still be considered valid. A hard constraint is a constraint that cannot be violated in a valid schedule. Any schedule that violates a hard constraint would be considered invalid.

Soft and hard constraints are implemented by varying the amount of penalty associated with each constraint. A hard constraint is given a higher penalty value than a soft constraint. This results in a schedule that violates hard constraints having a poorer fitness than a schedule that violates soft constraints. Since the scheduling algorithm favors those schedules that have a better fitness, schedules that violate hard constraints tend to die off while schedules that violate soft constraints tend to continue to be selected for reproduction.

The course scheduling system does not arbitrarily decide which constraints to enforce as hard constraints and which constraints to enforce as soft constraints, but rather gives this control to the scheduler. The scheduler can set the penalty value for each constraint, those that receive higher penalty values will be seen as hard constraints while those that receive lower penalty values will be seen as soft constraints.

3.3.9 Programming Language

The scheduling algorithm was implemented using the Microsoft Visual C++ programming language. Database access was performed using the Microsoft Foundation Classes Open Database Connectivity objects.

4. EVALUATION AND RESULTS

4.1 Testing

The modules within the scheduling algorithm and the GUI were tested as they were developed. When a module was completed it was tested with test cases chosen to exercise the paths through the module. If errors were found, they were corrected and the module was re-tested.

Testing for the completed system was performed by creating multiple schedules for a set of inputs and verifying the correctness of the schedules that were created. All constraints were verified by hand to ensure that they had been eliminated by the scheduling algorithm. In cases where a completed schedule contained constraint violations, it was verified that all constraint violations were reported by the scheduling algorithm. Non constraint issues such as verifying that day, time and room assignments were correct were also verified by hand.

Since the system is designed for non-technical users, the GUI is an extremely important feature. Because of this, the GUI's usability was tested by non-technical users. Each user was asked to work with the GUI and answer the following questions:

1. Were the forms/reports easy to use?
2. Were the required inputs easy to understand?
3. Were the operations uniform throughout the GUI?
4. How easy was it to learn to use the system? (scale of 1 (hard) to 5 (easy))
5. What is the best aspect of the GUI?
6. What is the poorest aspect of the GUI?

The responses to each of these questions have been considered and where appropriate have been used to correct problems and improve the usability of the GUI.

4.2 Evaluation

4.2.1 Running Time

The time required to compare constraints has the largest effect on the running time of the scheduling algorithm. The use of arrays to represent both schedules and constraints necessitates the use of nested loops to compare a single schedule to a single constraint. Each of the constraint comparing functions has two levels of nesting which gives them an upper bound running time of $O(n^2)$. Since the scheduling algorithm implements sGA, all schedules must have their fitness calculated in each generation. The *CalculateFitness* module performs this operation by looping through each schedule and invoking the various fitness functions. Using a loop to invoke functions that have two levels of nesting causes *CalculateFitness* to have three levels of nesting which gives it an upper bound running time of $O(n^3)$. The *CreateSchedule* module uses a loop to control the number of generations, and for each generation invokes the *CalculateFitness* module, which means that four levels of nested loops exist for *CreateSchedule* and gives an overall upper bound of $O(n^4)$ for the scheduling algorithm.

4.2.2 Violation Free Schedules

The scheduling algorithm is not always able to create schedules that do not violate any of the specified constraints. This is mainly due to the number of constraints that have been specified in the test cases (fewer constraints make it easier to find schedules that do not violate any constraints). However, the algorithm has been consistently able to

produce schedules that violate only a few constraints. Typically 2 to 5 constraint violations exist in a completed schedule.

This small number of violations can easily be manually corrected by the scheduler. To assist in this process the scheduling algorithm records all constraint violations that exist in a completed schedule, and also provides for hard and soft constraints as discussed in section 3.

4.2.3 Comparison of Simple GA and Steady State GA

This project implemented both sGA and ssGA for the purpose of determining which method of selection and replacement was most suited to the creation of course schedules. While the scheduler is interested in how long it takes to create a schedule, using time as a measure for comparison may not be completely accurate. Times may vary for reasons that do not involve the scheduling algorithm, such as the CPU's switching between multiple processes. A generation, as a unit of measure, provides a more accurate basis for comparison because a generation is an equal amount of work in both types of genetic algorithms. For these reasons generations will be used as the unit of comparison. Figure 4.1 compares the performance of sGA and ssGA over the first 1000 generations. For all comparisons an average for the violations was taken based on 10 separate runs of each GA type. Each chart used a separate run of 10 for its data.

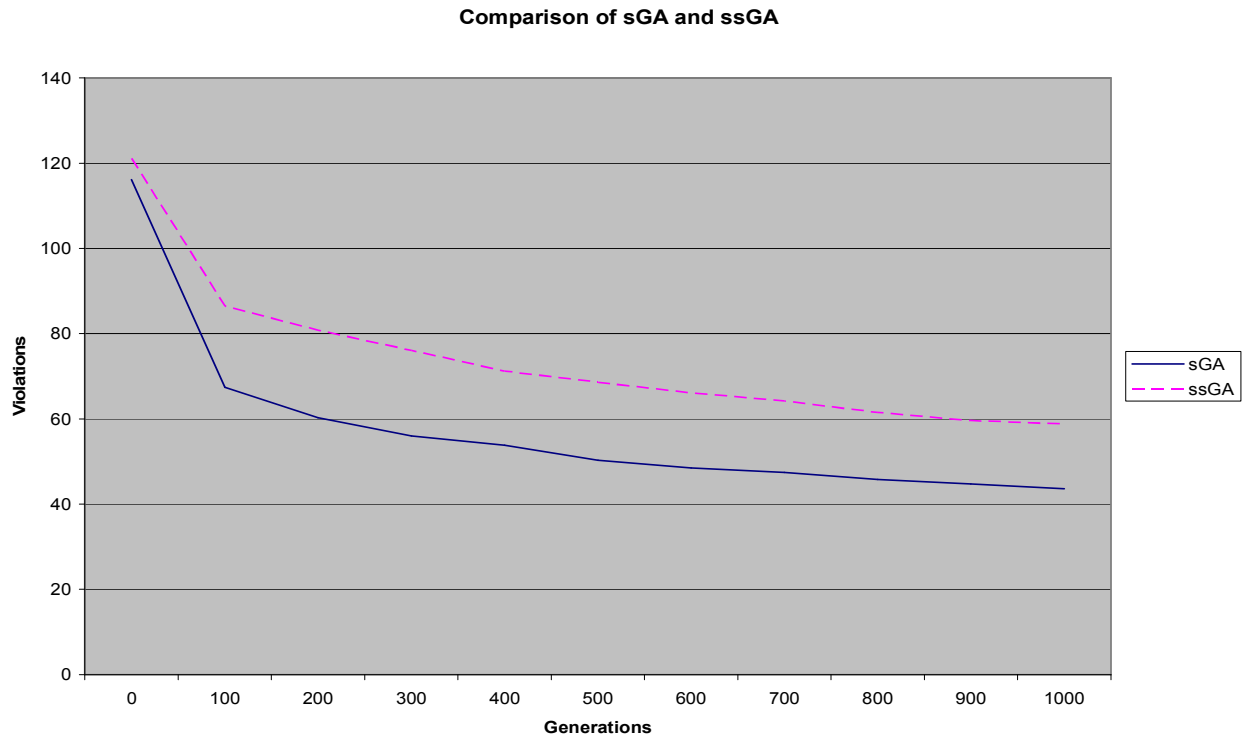


Figure 4.1 Comparison of sGA and ssGA Over 1000 Generations

Figure 4.2 compares the performance of sGA and ssGA over the first 10,000 generations.

The performance of both types of GA over the first 1000 generations is consistent with that seen in figure 4.1.




Figure 4.2 Comparison of sGA and ssGA Over 10,000 Generations

Figure 4.3 compares the performance of sGA and ssGA over 100,000 generations.

As seen in figures 4.1, 4.2, and 4.3 sGA outperforms ssGA from start to finish. Because of this, sGA is implemented in the final version of this project. Although the number of generations was used to make the comparison between sGA and ssGA, the scheduler will be interested in the amount of time it takes to create a schedule. Each of the schedules that have been created during the testing phase of this project has been created in 5 to 10 minutes on a desktop computer with a 2 GHz processor and 256 Mb of RAM. This is significant, because the schedules are being created very quickly on a computer that would be considered low end by today's standards.

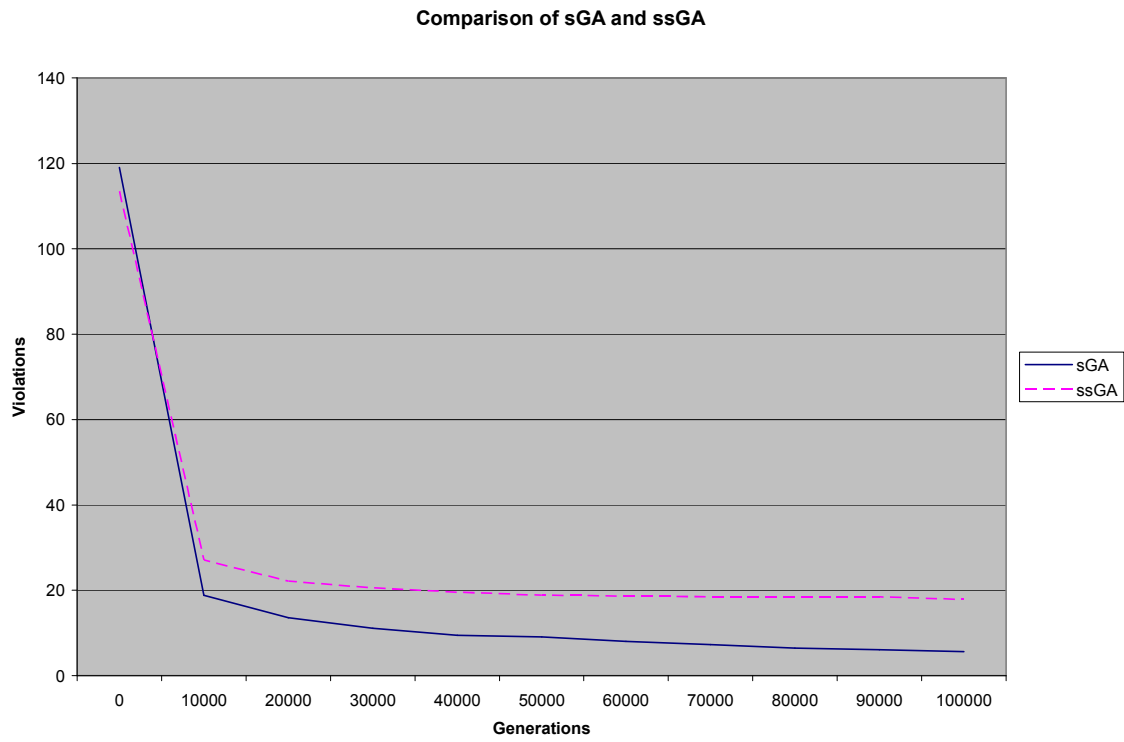


Figure 4.3 Comparison of sGA and ssGA Over 100,000 Generations

5. FUTURE WORK

5.1 Mechanism for Manual Changes

Since the scheduling algorithm is not always able to create a schedule that is free of constraint violations, a method could be made available to allow the scheduler to make manual changes within the GUI to a schedule to correct the constraint violations that still exist. As the scheduler makes a change to the schedule, the GUI could report any constraint violations caused by the change. The GUI could track all constraints that exist while the schedule is being altered and inform the scheduler when no constraint violations exist.

5.2 Determination of Data Dependencies

As discussed in section 3.2.4, the data entry by the scheduler is critical since conflicting constraint values can easily be entered. The relationships between each constraint type could be re-examined to determine if any conflicts have been missed and where possible add code that can prevent these conflicts from being input.

5.3 Testing On Larger Schedules

The scheduling algorithm has been shown to produce good schedules when the number of course sections to schedule is small (38 course sections were used in most tests). The scheduling algorithm could be tested using a larger number of course sections to determine if it can produce schedules with few constraint violations and measure the amount of time it takes to do so.

5.4 Use of Clusters in Schedule Creation

Originally proposed was the idea of creating a scheduling algorithm that would run on a cluster of nodes. Because of the success of the algorithm running on a single computer, it was determined that this step was not necessary. If in the future, it is determined that the scheduling algorithm does not perform well with a large number of course sections, then a clustered implementation can be pursued.

Each node in the cluster would run the same copy of the scheduling algorithm. Since genetic algorithms work at random, each copy of the genetic algorithm would arrive at different schedules. Each node would go through the process of creating new generations which would result in each node having a unique generation of schedules. At some regular interval, each node would choose the best schedules of the current generation to distribute to all other nodes. Each node would then have the same generation made up of the best schedules from each node. At this point, each node would continue to work in isolation until the next exchange occurs. The idea here is to determine if having the scheduling algorithm running on multiple nodes, with each node reaching its own random conclusions (and sharing those conclusions with all other nodes), will result in schedules that are better than those created on a single computer.

6. CONCLUSION

This paper presented a genetic algorithm that was designed and implemented to solve the course scheduling problem. The algorithm, based on user input, creates a semester course schedule by assigning each course a room, instructor, and meeting time. The scheduling algorithm considers 10 different constraints and attempts to create a schedule that does not violate any of these constraints.

The course scheduling system was designed to be used by non-technical users, so a complete graphical user interface was provided. The implementation of the GUI focused on creating a GUI that is easy to use and that is uniform in its operations.

Course schedules are created by applying the genetic operators (selection, recombination, and mutation) to the course scheduling problem. Two types of genetic algorithms (simple GA and steady state GA) were implemented and compared for effectiveness. Based on this comparison, simple GA was found to be the most effective method. For this reason the final product implements simple GA.

The scheduling algorithm implemented by this project has been found to be a very good solution to the course scheduling problem. Although the scheduling algorithm does not usually create violation free schedules, it does create schedules that have very few constraint violations (usually 2 to 5). The scheduling algorithm produces usable schedules that free the scheduler from the arduous task of manual schedule creation.

BIBLIOGRAPHY AND REFERENCES

- [Cormen 1990] Cormen, T., Leiserson, C., Rivest, R. *Introduction to Algorithms*, McGraw-Hill, New York, 1990.
- [Corne 1995] Corne, D., Ross, P. and Fang, H. *Evolving Timetables* In *Practical Handbook of Genetic Algorithms*, CRC Press, Boca Raton, 1995. Edited by Chambers, L.
- [Goldberg 1986] Goldberg, D. E. *The Genetic Algorithm: Who, How, and What Next?* In *Adaptive and Learning Systems*, Plenum, New York, 1986. Edited By Narendra, K.
- [Goldberg 1989] Goldberg, D. E. *Genetic Algorithms in Search, Optimization, and Machine Learning*, Addison-Wesley, New York, 1989.
- [Holland 1992] Holland, J. H., Genetic Algorithms, *Scientific American* 267(1) 1992, 66-72
- [Kragelund 1997] Kragelund, L. V. Solving a Timetabling Problem using Hybrid Genetic Algorithms. *Software – Practice And Experience* 27, (Oct. 1997), 1121-1134.
- [Pacheco 1997] Pacheco, P. S. *Parallel Programming With MPI*, Morgan Kaufmann, San Francisco, 1997.
- [Pfister 1998] Pfister, G. *In Search of Clusters*. Second Edition, Prentice-Hall PTR, UpperSaddle, NJ, 1998.
- [Sullivan 1997] Sullivan, M. Genetic Algorithms: A Brief Introduction. Available from http://www.cs.qub.ac.uk/~M.Sullivan/ga/ga_index.html (visited Sept. 1, 2002)
- [Wayner 1991] Wayner, P. Genetic Algorithms: Programming takes a valuable tip from nature, *Byte*, January, 1991, 361-368.

APPENDIX A – REPORTS

APPENDIX B – DATA DICTIONARY

Entity	Field	Data type	Length	Description
Buildings				
	BuildingId	Integer		System assigned unique value
	Name	String	20	Name of the building
Blackouts				
	Days	Integer		Integer value representing the days of the week for which the record exist.
	StartTime	Integer		Beginning time for the blackout – entered as military time
	EndTime	Integer		Ending time for the blackout – entered as military time
Classrooms				
	RoomId	Integer		System assigned unique value
	RoomNumber	String	6	Room number as assigned by university
	BuildingId	Integer		Unique value from the Buildings entity
	Capacity	Integer		Max number of people that can be in the room
	Active	Boolean		Determines if the room can be included on a schedule.
Courses				
	CourseId	Integer		System assigned unique value
	Subject	String	4	4 letter code representing the course subject such as ‘ENGL’ for English
	CourseNumber	String	4	4 number code representing the course such as ‘1301’
	Title	String	30	Name of the course
	Schedule	Boolean		Indicates whether or not the course should be placed on the schedule
	ContactHours	Integer		Number of hours actually spent in the classroom
CourseConflicts				
	Course1Id	Integer		Unique value from the Courses entity
	Course2Id	Integer		Unique value from the Courses entity
	SemId	Integer		Semester for which the constraint exist

Entity	Field	Data type	Length	Description
CourseClassrooms				
	CourseId	Integer		Unique value from the Courses entity
	RoomId	Integer		Unique value from the Classrooms table
	SemId	Integer		Semester for which the constraint exist
CourseInstructors				
	CourseId	Integer		Unique value from the Courses entity
	InstructorId	Integer		Unique value from the Instructors entity
	SemId	Integer		Semester for which the constraint exist
CourseTimes				
	CourseId	Integer		Unique value from the Courses entity
	Days	Integer		Integer value representing the days of the week on which the course should meet
	StartTime	Integer		Starting time for the course expressed in military time
	SemId	Integer		Semester for which the constraint exist
Instructors				
	InstructorId	Integer		System assigned unique value
	LastName	String	20	Instuctor's last name
	FirstName	String	20	Instructor's first name
InstructorCourses				
	InstructorId	Integer		Unique value from the Instructor entity
	CourseId	Integer		Unique value from the Courses entity
InstructorLoads				
	InstructorId	Integer		Unique value from the Instructor entity
	MinHours	Integer		Min number of hours the instructor will spend in the classroom for one week
	MaxHours	Integer		Max number of hours the instructor will spend in the classroom for one week
	SemId	Integer		Semester for which the constraint exist
InstructorPreferences				
	InstructorId	Integer		Unique value from the Instructor entity
	Days	Integer		Integer value representing the days the instructor prefers to teach
	StartTime	Integer		Starting time instructor prefers expressed in military time
	EndTime	Integer		Ending time instructor prefers expressed in military time.
	SemId	Integer		Semester for which the constraint exist

Entity	Field	Data type	Length	Description
Semesters				
	SemesterId	Integer		System assigned unique value
	SemesterName	String	10	Name of the semester. For example 'Fall'
	Year	String	4	Year of the semester. For example '2002'
Schedule				
	ScheduleId	Integer		System assigned unique value
	CrsSchedId	Integer		Courses to schedule primary key
	CourseId	Integer		Unique value from the Courses entity
	InstructorId	Integer		Unique value from the Instructor entity
	RoomId	Integer		Unique value from the Classrooms entity
	Days	Integer		Integer value representing the days of the week the course will meet
	StartTime	Integer		Starting time of the course expressed in military time
	EndTime	Integer		Ending time of the course expressed in military time
	SemesterId	Integer		Unique value from the Semesters entity. Indicates in which semester the course will be taught.
CoursesToSchedule				
	CrsSchedId	Integer		System generated primary key
	CourseId	Integer		Primary key of courses table
	SemId	Integer		Semester for which the course is to be scheduled
	Capacity	Integer		Max capacity for the course offering
ConstraintPenalties				
	CrsInst	Integer		Penalty value for course instructor constraint
	CrsCon	Integer		Penalty value for course conflict constraint
	CrsClass	Integer		Penalty value for course classroom constraint
	CrsTimes	Integer		Penalty value for course times constraint
	InstCrs	Integer		Penalty value for instructor course constraint
	InstLoad	Integer		Penalty value for instructor load constraint
	InstPref	Integer		Penalty value for instructor preferences constraint

Entity	Field	Data type	Length	Description
ConstraintViolations				
	Key	Integer		System assigned unique value
	CrsSchedId1	Integer		First course involved in the violation
	CrsSchedId2	Integer		Second course involved in the violation
	InstId	Integer		Instructor involved in the violation
	RoomId	Integer		Room involved in the violation
	ViolationType	Integer		Type of violation
	ScheduleId	Integer		Schedule to which the violation belongs
ScheduleDates				
	ScheduleId	Integer		System assigned unique value
	Title	String	50	User defined title
	DateCreated	Date/Time		Date schedule was created
Keys				
	Instructors	Integer		Largest assigned primary key in the instructors table
	Buildings	Integer		Largest assigned primary key in the buildings table
	Classrooms	Integer		Largest assigned primary key in the classrooms table
	Schedule	Integer		Largest assigned primary key in the schedule table
	Courses ToSchedule	Integer		Largest assigned primary key in the CoursesToSchedule table
	KeysSaved	Boolean		Indicates whether the key values has been saved by the GUI